

RandomForest

April 18, 2021

```
[1]: # library preparations
import scipy.io
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import sklearn
import joblib
import seaborn as sns
import time
import os

# export PATH=/usr/local/texlive/2021/bin/x86_64-linux/xelatex:$PATH
# jupyter nbconvert RandomForest.ipynb --to pdf
print(os.environ['PATH'])
```

```
/home/samsoup/anaconda3/bin:/home/samsoup/anaconda3/condabin:/usr/local/sbin:/us
r/local/bin:/usr/sbin:/usr/bin:/sbin:/bin:/usr/games:/usr/local/games:/snap/bin
```

1 Model Fitting

Given that we could just ask the simulation for more samples, I decided to NOT employ a typical training-split. Rather, I will deploy all samples for training and optimize based on cross validation approaches, given every test record a turn at being in the training set and test set. Once the model is fitted, we can then ask the simulation for some more samples (say 1000) to use as a test set completely related here.

1.1 Model Evaluations

Many metrics can be used to evaluate models, some I calculate here are:

1. Accuracy: $TP + TN / \text{total}$, this is the number of samples the RF model classifies correctly
2. Error Rate: $1 - \text{Accuracy}$, this is the number of samples the RF model classifies incorrectly
3. True Positive Rate (TPR) | Sensitivity | Recall | Coverage: $TP / (TP + FN)$, fraction of SAS examples correctly predicted
 - There's typically a trade off between Recall and Precision below
4. True Negative Rate (TNR) | Specificity: $TN / (FP + TN)$, fraction of neutral examples correctly predicted
5. False Positive Rate (FPR): $FP / (FP + TN)$ fraction of neutral examples predicted as having SAS (really bad)

6. False Negative Rate (FNR): $\text{FN} / (\text{TP} + \text{FN})$ fraction of SAS examples predicted as having SAS (not as bad, but still bad)
7. Precision: $\text{TP} / (\text{TP} + \text{FP})$, fraction of samples that actually have SAS out of total samples predicted to have SAS
 - Precision addresses the question: “Given a sample predicted to have SAS, how likely is it to be correct?”
 - **We may want to sacrifice Recall in order to achieve a high precision**
8. F-measure: $\frac{2 * \text{precision} * \text{recall}}{\text{precision} + \text{recall}}$, a harmonic mean (meaning the resulting metric is *closer to the smaller of the input, that is, F-measure is closer to either precision or recall, whichever is smaller in magnitude) of precision and recall
 - **Ideally, a F-measure should be high and indicate that both precision and recall are high**
9. Area under the Curve (AUC): this is the area of the curve under the Receiver operating characteristic (ROC) curve, which demonstrate the specificity-sensitivity trade-off (TPR vs TNR)

1.2 Confusion Matrix

1.3 Alternative Metric: Cost Matrix

Sometimes, in the realm of health sciences, we want to *punish* or *reward* the model for doing something really well compared to other. For example, in cancer prediction, more emphasis is placed on avoiding False Negatives (not detecting the cancer), so that we may wish to assign costs/weights (negative means reward) to TP, FP, TN, and FN like this:

This can be implemented as other alternative to above metrics during cross validation. Or, cost matrix can be used to classify one particular record. That is, **we can use cost matrix to evaluate risk**

With a RandomForest, I am able to extract the *probability* of a sample as showing SAS or not, then say:

$$P(\text{SAS}) = 0.2$$

$$P(\text{neutral, other}) = 0.8$$

Given the above cost matrix, then when I:

- Classify this sample as having SAS, there's a 0.8 probability I am wrong, making a FP error with cost 1, or a risk of 0.8
- Classify this sample as not having SAS, there's a 0.2 probability I am wrong, making a FN error with cost 100, or a risk of 20
- I can then decide to classify this sample based on probability or this new risk factor

```
[2]: load_train = np.load('./data/train.npz', allow_pickle=True)
load_test = np.load('./data/test.npz', allow_pickle=True)
X_train, y_train = load_train['X_train'], load_train['y_train']
X_test, y_test = load_test['X_test'], load_test['y_test']

samples, rows, cols = X_train.shape
print("Before flattening:")
print(X_train.shape)
```

```
print(y_train.shape)
print(X_test.shape)
print(y_test.shape)
```

Before flattening:

```
(10000, 160, 2)
(10000, 1)
(1000, 160, 2)
(1000, 1)
```

[3] :

```
"""
fst should range anywhere between 0 and 1 in practice,
and mean squared error is in the same bounds.
```

*But whether we want equal sized bins or not is still not clear
(we may want higher granularity for bins with high fst and good fit,
and low for low fst regardless of fit)*

```
"""
```

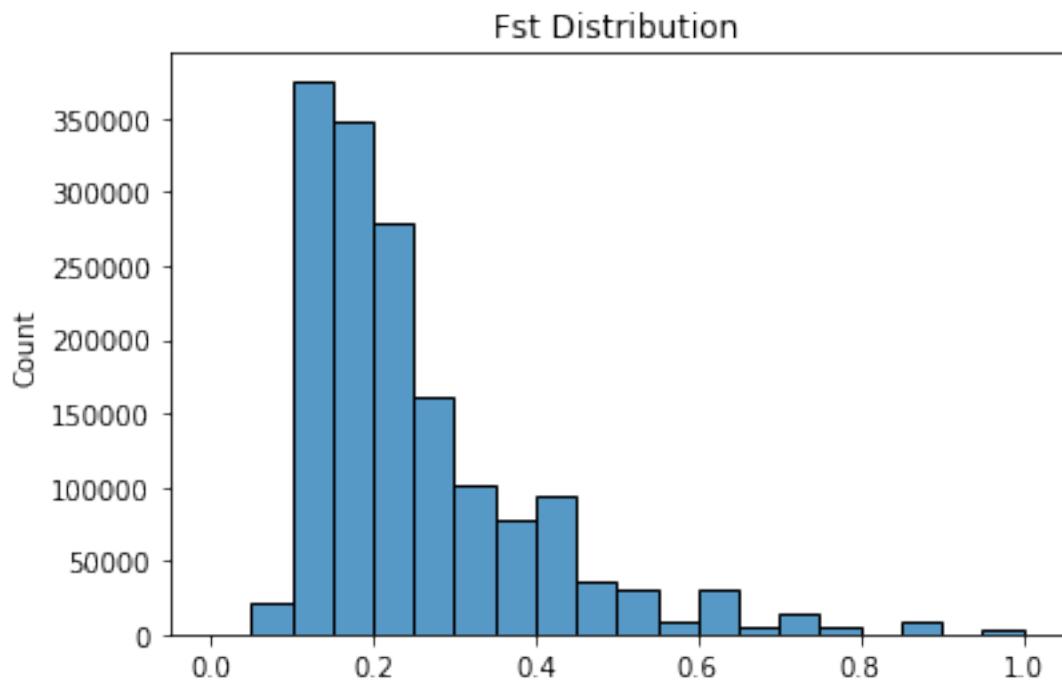
```
# check ranges for f_st
print("Min Fst", np.amin(X_train[:, :, 0]))
print("Max Fst", np.amax(X_train[:, :, 0]))

# check ranges for means squared error
print("Min SSE", np.amin(X_train[:, :, 1]))
print("Max SSE", np.amax(X_train[:, :, 1]))
```

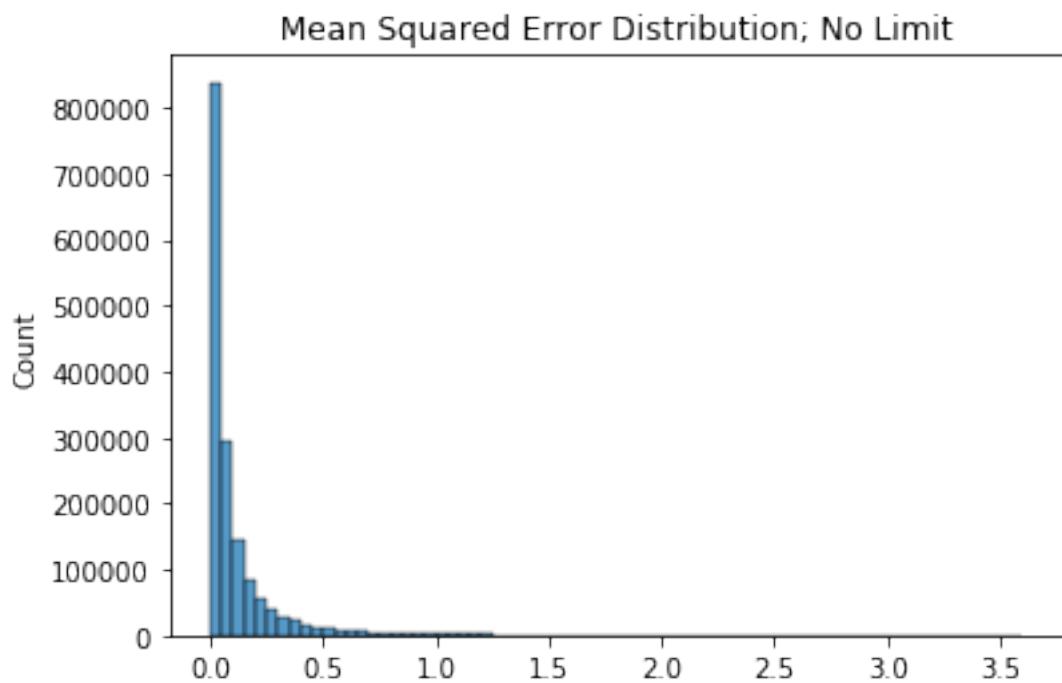
```
Min Fst 0.0
Max Fst 1.0
Min SSE 0.0
Max SSE 3.5890404752572502
```

[4] :

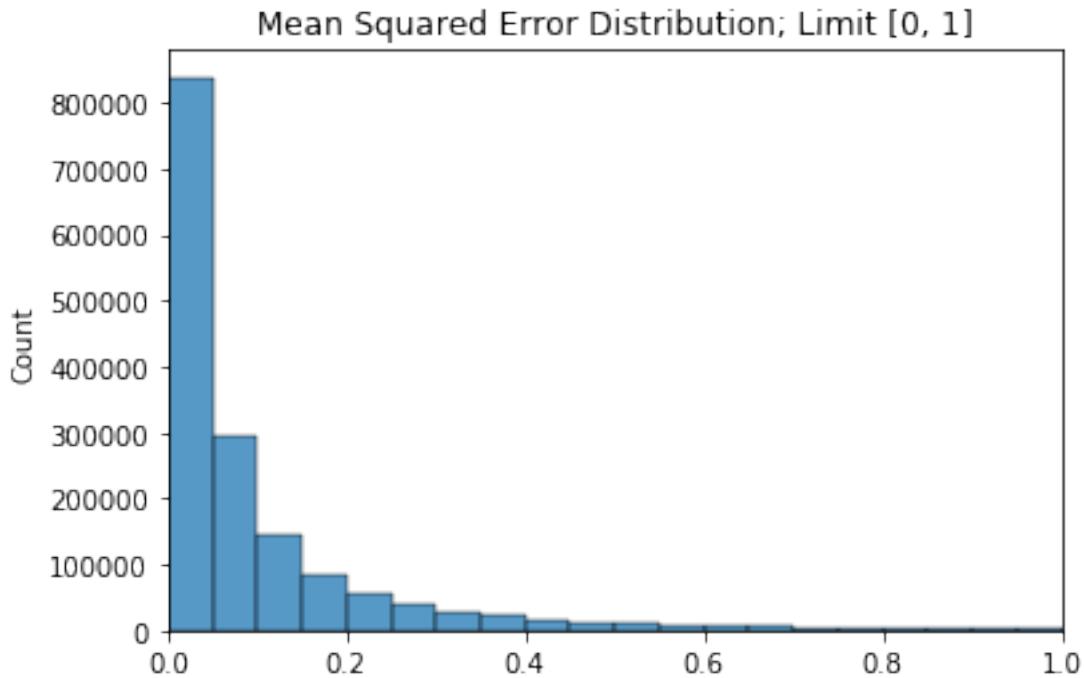
```
sns.histplot(X_train[:, :, 0].reshape(-1), binwidth=0.05)
plt.title("Fst Distribution")
plt.show()
```



```
[5]: sns.histplot(X_train[:, :, 1].reshape(-1), binwidth=0.05)
plt.title("Mean Squared Error Distribution; No Limit")
plt.show()
```



```
[6]: sns.histplot(X_train[:, :, 1].reshape(-1), binwidth=0.05)
plt.title("Mean Squared Error Distribution; Limit [0, 1]")
plt.xlim([0, 1])
plt.show()
```



```
[7]: # ~ 1% of metrics have MSE >= 1,
print(sum(X_train[:, :, 1].reshape(-1) >= 1))

# ~3% of metrics have MSE >= 0.6
print(sum(X_train[:, :, 1].reshape(-1) >= 0.6))

# ~6% of metrics have MSE >= 0.4
print(sum(X_train[:, :, 1].reshape(-1) >= 0.4))
```

15905
49943
98852

```
[8]: """
Based on the plots, let's make 5 bins each for Fst and MSE, namely:
    1           2           3           4           5
Fst: [0, 0.2), [0.2, 0.4), [0.4, 0.6), [0.6, 0.8), [0.8, 1.0+0.1)
MSE: [0, 0.1), [0.1, 0.2), [0.2, 0.3), [0.3, 0.4), [0.4, max+0.1)
```

This way, we abstract away the 'location' dimension as a factor

```

"""
Fst_bins = list(i / 100 for i in range(0, 100, 5))
# Fst_bins = [0, 0.4]
Fst_bins.append(np.amax(X_train[:, :, 0]) + 0.01)
print("Fst bins", Fst_bins)
MSE_bins = list(i / 10 for i in range(0, 5, 1))
# MSE_bins = [0, 0.4]
MSE_bins.append(np.amax(X_train[:, :, 1]) + 0.01)
print("MSE bins", MSE_bins)

def makeLabels(iterable):
    # takes an iterable of bin (start, stop) and make labels
    return ["-".join([str(start), str(round(stop, 2))]) for start, stop in
            zip(iterable, iterable[1:])]
Fst_labels = makeLabels(Fst_bins)
MSE_labels = makeLabels(MSE_bins)

print("Fst labels", Fst_labels)
print("MSE labels", MSE_labels)

def assignBins(Fst_bins, MSE_bins, X_train, X_test):
    Fst_assignments = np.digitize(X_train[:, :, 0], Fst_bins)
    Fst_bins_count = len(Fst_bins)

    MSE_assignments = np.digitize(X_train[:, :, 1], MSE_bins)
    MSE_bins_count = len(MSE_bins)

    # rows: Fst bin #, cols: MSE bin
    X_train_bins = np.zeros((Fst_bins_count-1, MSE_bins_count-1), dtype=np.int)
    X_test_bins = np.zeros((Fst_bins_count-1, MSE_bins_count-1), dtype=np.int)

    def computeMatrix(base, matrix):
        for i in range(0, len(matrix)):
            sample = matrix[i]
            rows = sample.shape[0]

            # transform this (160, 2) matrix into a fzm bin matrix
            empty = np.zeros((Fst_bins_count-1, MSE_bins_count-1), dtype=np.int)
            for j in range(0, rows):
                Fst_bin = Fst_assignments[i, j]
                MSE_bin = MSE_assignments[i, j]
                empty[Fst_bin-1, MSE_bin-1] += 1
            base = np.concatenate((base, empty))

        base = base.reshape(-1, Fst_bins_count-1, MSE_bins_count-1)

```

```

# drop the first entry (only used for appending)
base = base[1:, :, :]

return base

X_train_bins = computeMatrix(X_train_bins, X_train)
X_test_bins = computeMatrix(X_test_bins, X_test)

return X_train_bins, X_test_bins

X_train_bins, X_test_bins = assignBins(Fst_bins, MSE_bins, X_train, X_test)
print(X_train_bins.shape)
print(X_test_bins.shape)

```

```

Fst bins [0.0, 0.05, 0.1, 0.15, 0.2, 0.25, 0.3, 0.35, 0.4, 0.45, 0.5, 0.55, 0.6,
0.65, 0.7, 0.75, 0.8, 0.85, 0.9, 0.95, 1.01]
MSE bins [0.0, 0.1, 0.2, 0.3, 0.4, 3.59904047525725]
Fst labels ['0.0-0.05', '0.05-0.1', '0.1-0.15', '0.15-0.2', '0.2-0.25',
'0.25-0.3', '0.3-0.35', '0.35-0.4', '0.4-0.45', '0.45-0.5', '0.5-0.55',
'0.55-0.6', '0.6-0.65', '0.65-0.7', '0.7-0.75', '0.75-0.8', '0.8-0.85',
'0.85-0.9', '0.9-0.95', '0.95-1.01']
MSE labels ['0.0-0.1', '0.1-0.2', '0.2-0.3', '0.3-0.4', '0.4-3.6']
(10000, 20, 5)
(1000, 20, 5)

```

```

[9]: def generateHeatmaps(matrices, Fst_labels, MSE_labels, SAS_labels, savePath):
    for i in range(0, matrices.shape[0]):
        ax = sns.heatmap(
            pd.DataFrame(
                matrices[i, :, :],
                index=Fst_labels,
                columns=MSE_labels
            ),
            linewidths=.5)
        # invert y axis to sort Fst ascending
        ax.invert_yaxis()
        figure = plt.gcf()
        figure.set_size_inches(16, 12)
        plt.title("Output " + str(i+1) + " - " + SAS_labels[i] + " SNS")
        plt.ylabel("Fst")
        plt.xlabel("MSE")

        # save plot
        saveFile = os.path.join(savePath, "output"+str(i+1)+".png")
        plt.savefig(saveFile)
        plt.close() # do NOT show in-line

```

```

generateHeatmaps(
    # just do the first 10 to save some time
    X_train_bins[:10],
    Fst_labels, MSE_labels, y_train, savePath=".\\Images")

```

```
[10]: y_train = y_train.reshape(y_train.shape[0], )
y_test = y_test.reshape(y_test.shape[0], )
X_train_bins = X_train_bins.reshape(X_train_bins.shape[0], -1)
X_test_bins = X_test_bins.reshape(X_test.shape[0], -1)

print("After flattening:")
print(X_train_bins.shape)
print(y_train.shape)
print(X_test_bins.shape)
print(y_test.shape)
```

After flattening:

```
(10000, 100)
(10000,)
(1000, 100)
(1000,)
```

```
[11]: # Prepare for K fold cross validation
from sklearn.model_selection import train_test_split
from sklearn.model_selection import KFold
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score, f1_score
from sklearn.metrics import precision_score, recall_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import roc_auc_score

k = 10
# again, random_state for sake of reproducibility
kf = KFold(n_splits=k, random_state = 42, shuffle=True)

# default confusion matrix
# we can calculate TP, FP, FN, TN from this matrix at the end
matrix = pd.DataFrame(0, index=['True Yes', 'True No', ''],
                      columns=['Pred Yes', 'Pred No', ''])

start = time.time()
for train_index, test_index in kf.split(X_train_bins):
    # identify the train and test set within each fold
    X_fold_train, X_fold_test = X_train_bins[train_index], ↵
    ↵X_train_bins[test_index]
    y_fold_train, y_fold_test = y_train[train_index], y_train[test_index]
```

```

Time taken: 13.930890083312988
          Pred Yes  Pred No
True  Yes      2740      2089    4829
True  No       2260      2911    5171
                  5000      5000   10000

```

Accuracy: 0.5651
Error Rate: 0.4349
True Positive Rate (TPR) | Sensitivity | Recall | Coverage

```

0.5674052598881756
True Negative Rate (TNR) | Specificity
0.5629472055695224
False Positive Rate (FPR) is
0.4370527944304777
False Negative Rate (FNR) is
0.4325947401118244
Average Precision: 0.548
Average Recall:
0.5674052598881756
Average F Measures:
0.557533828466782

```

```
[12]: # https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.
      ↪RandomForestClassifier.html
# n_estimators = number of trees to bag
# Fit the model on all training samples possible
# clf = RandomForestClassifier(n_estimators=500, criterion='gini', ↪
#     ↪max_depth=None,
#                                     min_samples_split=2, min_samples_leaf=1, ↪
#     ↪min_weight_fraction_leaf=0.0,
#                                     max_features='auto', max_leaf_nodes=None, ↪
#     ↪min_impurity_decrease=0.0,
#                                     min_impurity_split=None, bootstrap=True, ↪
#     ↪oob_score=False, n_jobs=None,
#                                     random_state=None, verbose=0, warm_start=False, ↪
#     ↪class_weight=None,
#                                     ccp_alpha=0.0, max_samples=None)
# Tune: min_samples_split, max_leaf_nodes, max_depth and min_samples_leaf.
model = RandomForestClassifier(n_estimators=500, oob_score=True,
                                max_features = 'sqrt', n_jobs=8)
# model = RandomForestClassifier(n_estimators=500, oob_score=True,
#                                min_samples_leaf = 10, max_depth = 5,
#                                min_samples_split = 100, max_leaf_nodes = 50,
#                                max_features = 'sqrt', n_jobs=8)
start = time.time()
model.fit(X_train_bins, y_train)
end = time.time()
print('Time taken:', end - start)

# save the model
# joblib.dump(model, "./data/random_forest_10000.joblib")

from itertools import product
feature_labels = [
    "Fst " + i + "; MSE " + j for i, j in product(Fst_labels, MSE_labels)
]
```

Time taken: 1.4501984119415283

Feature Fst 0.0-0.05; MSE 0.0-0.1 importance: 5.6529528361442624e-05

Feature Fst 0.0-0.05; MSE 0.1-0.2 importance: 0.0

Feature Fst 0.0-0.05; MSE 0.2-0.3 importance: 0.0

Feature Fst 0.0-0.05; MSE 0.3-0.4 importance: 0.0

Feature Fst 0.0-0.05; MSE 0.4-3.6 importance: 0.0

Feature Fst 0.05-0.1; MSE 0.0-0.1 importance: 0.038965607250915565

Feature Fst 0.05-0.1; MSE 0.1-0.2 importance: 5.16413219930819e-05

Feature Fst 0.05-0.1; MSE 0.2-0.3 importance: 6.658509726375368e-06

Feature Fst 0.05-0.1; MSE 0.3-0.4 importance: 0.0

Feature Fst 0.05-0.1; MSE 0.4-3.6 importance: 0.0

Feature Fst 0.1-0.15; MSE 0.0-0.1 importance: 0.06989721072270899

Feature Fst 0.1-0.15; MSE 0.1-0.2 importance: 4.3716218940199974e-05

Feature Fst 0.1-0.15; MSE 0.2-0.3 importance: 4.431587265636601e-06

Feature Fst 0.1-0.15; MSE 0.3-0.4 importance: 0.0

Feature Fst 0.1-0.15; MSE 0.4-3.6 importance: 0.0

Feature Fst 0.15-0.2; MSE 0.0-0.1 importance: 0.07006934186467917

Feature Fst 0.15-0.2; MSE 0.1-0.2 importance: 0.001044109448531129

Feature Fst 0.15-0.2; MSE 0.2-0.3 importance: 0.0

Feature Fst 0.15-0.2; MSE 0.3-0.4 importance: 0.0

Feature Fst 0.15-0.2; MSE 0.4-3.6 importance: 0.0

Feature Fst 0.2-0.25; MSE 0.0-0.1 importance: 0.06909830437691546

Feature Fst 0.2-0.25; MSE 0.1-0.2 importance: 0.026406221337038606

Feature Fst 0.2-0.25; MSE 0.2-0.3 importance: 0.00013774434030870155

Feature Fst 0.2-0.25; MSE 0.3-0.4 importance: 0.0

Feature Fst 0.2-0.25; MSE 0.4-3.6 importance: 0.0

Feature Fst 0.25-0.3; MSE 0.0-0.1 importance: 0.05850197698623722

Feature Fst 0.25-0.3; MSE 0.1-0.2 importance: 0.052862985642808925

Feature Fst 0.25-0.3; MSE 0.2-0.3 importance: 0.004842066071510723

Feature Fst 0.25-0.3; MSE 0.3-0.4 importance: 7.197749773550829e-05

Feature Fst 0.25-0.3; MSE 0.4-3.6 importance: 0.0

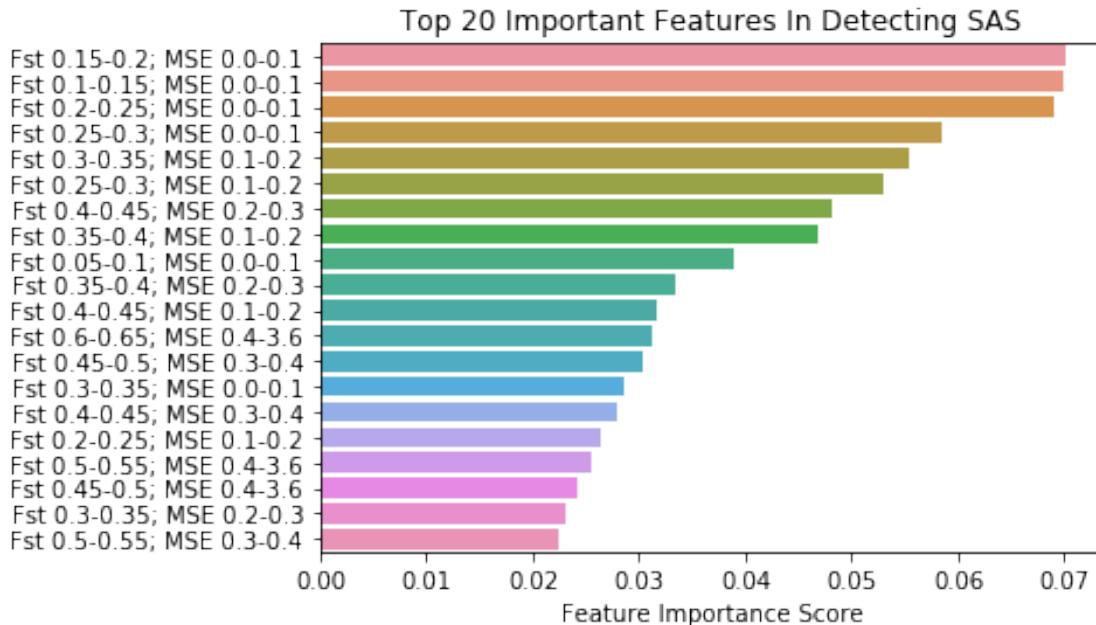
Feature Fst 0.3-0.35; MSE 0.0-0.1 importance: 0.028579186297277865

Feature Fst 0.3-0.35; MSE 0.1-0.2 importance: 0.05543667996238436
Feature Fst 0.3-0.35; MSE 0.2-0.3 importance: 0.023183360022508377
Feature Fst 0.3-0.35; MSE 0.3-0.4 importance: 0.0025230466569873154
Feature Fst 0.3-0.35; MSE 0.4-3.6 importance: 0.00011062565290891459
Feature Fst 0.35-0.4; MSE 0.0-0.1 importance: 0.011444215817121371
Feature Fst 0.35-0.4; MSE 0.1-0.2 importance: 0.04679918754601216
Feature Fst 0.35-0.4; MSE 0.2-0.3 importance: 0.03351890379935994
Feature Fst 0.35-0.4; MSE 0.3-0.4 importance: 0.008415574467479441
Feature Fst 0.35-0.4; MSE 0.4-3.6 importance: 0.0011874831912088055
Feature Fst 0.4-0.45; MSE 0.0-0.1 importance: 0.0028988808271889015
Feature Fst 0.4-0.45; MSE 0.1-0.2 importance: 0.03174826843871468
Feature Fst 0.4-0.45; MSE 0.2-0.3 importance: 0.048106314004522116
Feature Fst 0.4-0.45; MSE 0.3-0.4 importance: 0.027944283671902853
Feature Fst 0.4-0.45; MSE 0.4-3.6 importance: 0.010681271846408767
Feature Fst 0.45-0.5; MSE 0.0-0.1 importance: 0.00024393026879368152
Feature Fst 0.45-0.5; MSE 0.1-0.2 importance: 0.003569771655125391
Feature Fst 0.45-0.5; MSE 0.2-0.3 importance: 0.019862070209666374
Feature Fst 0.45-0.5; MSE 0.3-0.4 importance: 0.030410550071353352
Feature Fst 0.45-0.5; MSE 0.4-3.6 importance: 0.024102740441868373
Feature Fst 0.5-0.55; MSE 0.0-0.1 importance: 4.6080394153498264e-05
Feature Fst 0.5-0.55; MSE 0.1-0.2 importance: 0.001487794354915026
Feature Fst 0.5-0.55; MSE 0.2-0.3 importance: 0.010067306258679383
Feature Fst 0.5-0.55; MSE 0.3-0.4 importance: 0.022425812645846154
Feature Fst 0.5-0.55; MSE 0.4-3.6 importance: 0.025564403693927158
Feature Fst 0.55-0.6; MSE 0.0-0.1 importance: 1.4334265367203021e-05
Feature Fst 0.55-0.6; MSE 0.1-0.2 importance: 0.0001310509657601731
Feature Fst 0.55-0.6; MSE 0.2-0.3 importance: 0.001445052496849532
Feature Fst 0.55-0.6; MSE 0.3-0.4 importance: 0.007460453542753818
Feature Fst 0.55-0.6; MSE 0.4-3.6 importance: 0.019153073282765866
Feature Fst 0.6-0.65; MSE 0.0-0.1 importance: 9.571471871938711e-06
Feature Fst 0.6-0.65; MSE 0.1-0.2 importance: 0.00011173311469515142
Feature Fst 0.6-0.65; MSE 0.2-0.3 importance: 0.001014278857841836
Feature Fst 0.6-0.65; MSE 0.3-0.4 importance: 0.005537028350516485
Feature Fst 0.6-0.65; MSE 0.4-3.6 importance: 0.03123356216632302
Feature Fst 0.65-0.7; MSE 0.0-0.1 importance: 0.0
Feature Fst 0.65-0.7; MSE 0.1-0.2 importance: 1.7987227690951162e-05
Feature Fst 0.65-0.7; MSE 0.2-0.3 importance: 6.609691543692472e-05
Feature Fst 0.65-0.7; MSE 0.3-0.4 importance: 0.0007526370130214901
Feature Fst 0.65-0.7; MSE 0.4-3.6 importance: 0.015303125025956453
Feature Fst 0.7-0.75; MSE 0.0-0.1 importance: 0.0
Feature Fst 0.7-0.75; MSE 0.1-0.2 importance: 4.845780093890819e-06
Feature Fst 0.7-0.75; MSE 0.2-0.3 importance: 2.009394553183249e-05
Feature Fst 0.7-0.75; MSE 0.3-0.4 importance: 0.0002074271301588076
Feature Fst 0.7-0.75; MSE 0.4-3.6 importance: 0.02083266583603696
Feature Fst 0.75-0.8; MSE 0.0-0.1 importance: 0.0
Feature Fst 0.75-0.8; MSE 0.1-0.2 importance: 0.0
Feature Fst 0.75-0.8; MSE 0.2-0.3 importance: 0.0
Feature Fst 0.75-0.8; MSE 0.3-0.4 importance: 1.2023830178896957e-05

```

Feature Fst 0.75-0.8; MSE 0.4-3.6 importance: 0.012744495635065438
Feature Fst 0.8-0.85; MSE 0.0-0.1 importance: 0.0
Feature Fst 0.8-0.85; MSE 0.1-0.2 importance: 0.0
Feature Fst 0.8-0.85; MSE 0.2-0.3 importance: 0.0
Feature Fst 0.8-0.85; MSE 0.3-0.4 importance: 0.0
Feature Fst 0.8-0.85; MSE 0.4-3.6 importance: 0.0
Feature Fst 0.85-0.9; MSE 0.0-0.1 importance: 0.0
Feature Fst 0.85-0.9; MSE 0.1-0.2 importance: 0.0
Feature Fst 0.85-0.9; MSE 0.2-0.3 importance: 0.0
Feature Fst 0.85-0.9; MSE 0.3-0.4 importance: 0.0
Feature Fst 0.85-0.9; MSE 0.4-3.6 importance: 0.015316146321871636
Feature Fst 0.9-0.95; MSE 0.0-0.1 importance: 0.0
Feature Fst 0.9-0.95; MSE 0.1-0.2 importance: 0.0
Feature Fst 0.9-0.95; MSE 0.2-0.3 importance: 0.0
Feature Fst 0.9-0.95; MSE 0.3-0.4 importance: 0.0
Feature Fst 0.9-0.95; MSE 0.4-3.6 importance: 0.0
Feature Fst 0.95-1.01; MSE 0.0-0.1 importance: 0.0
Feature Fst 0.95-1.01; MSE 0.1-0.2 importance: 0.0
Feature Fst 0.95-1.01; MSE 0.2-0.3 importance: 0.0
Feature Fst 0.95-1.01; MSE 0.3-0.4 importance: 0.0
Feature Fst 0.95-1.01; MSE 0.4-3.6 importance: 0.00619405192222075

```



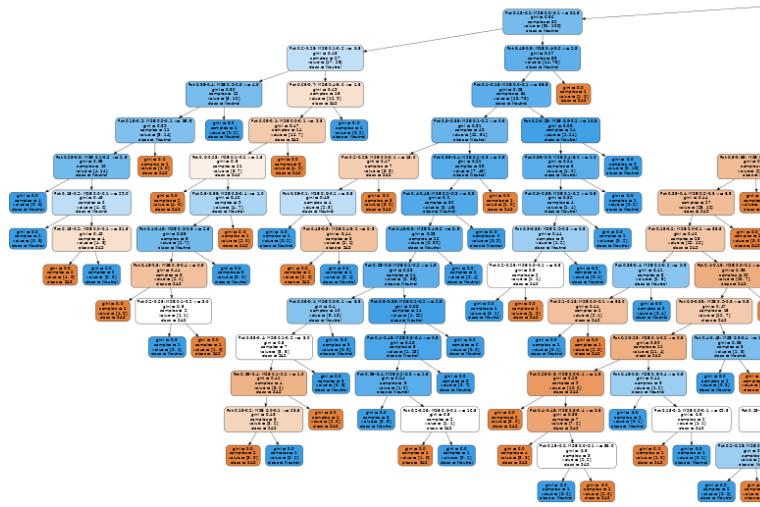
```
[13]: # Extract single tree
estimator = model.estimators_[203]

from sklearn.tree import export_graphviz
```

```
# Export as dot file
export_graphviz(estimator, out_file='tree.dot',
                feature_names = feature_labels,
                class_names = ["SAS", "Neutral"],
                rounded = True, proportion = False,
                precision = 2, filled = True)

# Convert to png using system command (requires Graphviz)
from subprocess import call
call(['dot', '-Tpng', 'tree.dot', '-o', 'tree.png', '-Gdpi=600'])

# Display in jupyter notebook
from IPython.display import Image
display(Image(filename = 'tree.png'))
```



```
[14]: # from sklearn.ensemble import RandomForestClassifier
# from sklearn.preprocessing import StandardScaler
# from sklearn.model_selection import cross_val_score
# from sklearn.pipeline import Pipeline
# from sklearn.model_selection import cross_val_predict
# from sklearn.metrics import confusion_matrix
# from sklearn.metrics import classification_report
# from sklearn.model_selection import GridSearchCV
# from sklearn.metrics import accuracy_score
# from sklearn.preprocessing import MinMaxScaler

# param_grid = {
#     'max_depth':list(range(35,55+1)),
#     'min_samples_leaf':list(range(8,35,2)),
#     'max_features':['sqrt', 'log2'],
# }

# clf = GridSearchCV(RandomForestClassifier(criterion='gini'), param_grid,
#                     scoring="accuracy", cv=5, n_jobs=8)
# clf.fit(X_train_bins, y_train)
# print(clf.best_params_)

# rf_pred = cross_val_predict(
#     RandomForestClassifier(criterion='gini'),
#     X=X_train_bins, y=y_train, cv=5, n_jobs=12)
# print("The average nested cross validation accuracy is ",accuracy_score(y_train, rf_pred))

# print(confusion_matrix(y_train, rf_pred))

# print(classification_report(y_train, rf_pred))
```

1.4 Test Set Evaluation

Below are performance measures on the test set never touched during model building

```
[15]: y_pred = model.predict(X_test_bins)
print("Accuracy on Test Set:", accuracy_score(y_test,y_pred))
print(f'Our OOB prediction of accuracy is: {model.oob_score_ * 100}%')

# show stats
matrix = pd.crosstab(y_pred, y_test, margins=True)
matrix.index = ['True Yes', 'True No', '']
matrix.columns = ['Pred Yes', 'Pred No', '']
print(matrix)
```

```

TP, FN = matrix.iloc[0, 0], matrix.iloc[0, 1]
FP, TN = matrix.iloc[1, 0], matrix.iloc[1, 1]
total = TP + FN + FP + TN
print("Accuracy: \t\t\t\t\t\t\t", (TP + TN) / total)
print("Error Rate: \t\t\t\t\t\t\t", (FN + FP) / total)
print("True Positive Rate (TPR) | Sensitivity | Recall | Coverage\t",
      TP / (TP + FN))
print("True Negative Rate (TNR) | Specificity \t\t\t\t\t",
      TN / (FP + TN))
print("False Positive Rate (FPR) is \t\t\t\t", FP / (FP + TN))
print("False Negative Rate (FNR) is \t\t\t\t", FN / (TP + FN))
precision = TP / (TP + FP)
recall = TP / (TP + FN)
print("Average Precision: \t\t\t\t", precision)
print("Average F Measures: \t\t\t\t",
      (2*recall*precision) / (recall + precision))
pred_prob = model.predict_proba(X_test_bins)[:, 1]
print("Average AUC: \t\t\t\t\t\t\t", roc_auc_score(y_test, pred_prob))

```

Accuracy on Test Set: 0.489
Our OOB prediction of accuracy is: 56.8%

	Pred Yes	Pred No
True Yes	238	249
True No	262	513
	500	500
		1000

Accuracy: 0.489
Error Rate: 0.511
True Positive Rate (TPR) | Sensitivity | Recall | Coverage
0.4887063655030801
True Negative Rate (TNR) | Specificity
0.48927875243664715
False Positive Rate (FPR) is
0.5107212475633528
False Negative Rate (FNR) is
0.5112936344969199
Average Precision: 0.476
Average F Measures:
0.4822695035460993
Average AUC:
0.5153840000000001

[16]: `from sklearn.metrics import precision_recall_curve`

```

# precision recall curve for training set
y_prob = model.predict_proba(X_test_bins)[:,1]
precisions, recalls, thresholds = precision_recall_curve(

```

```

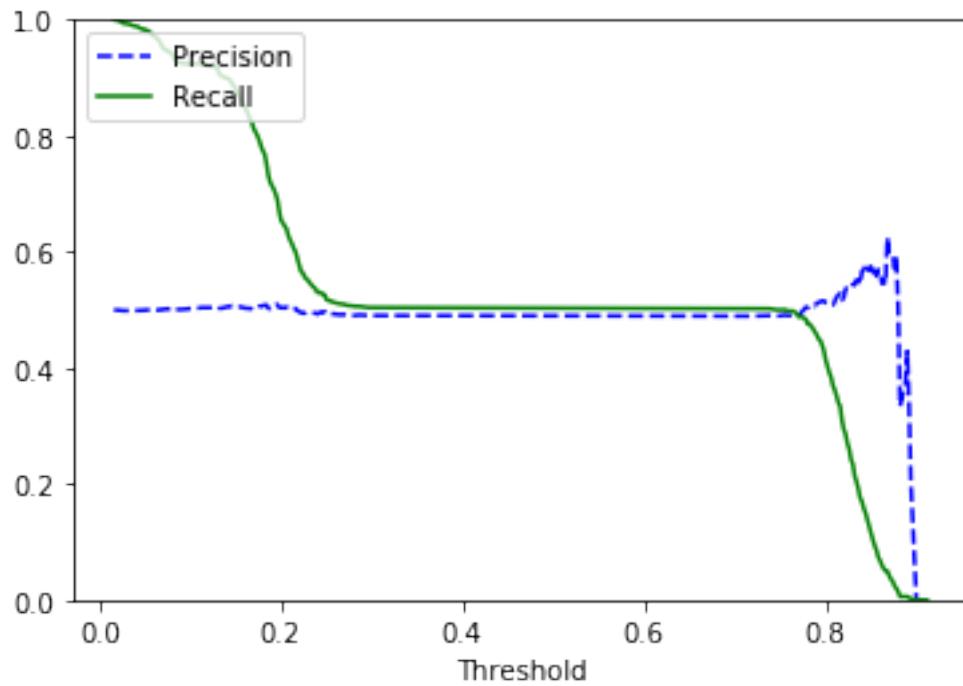
y_test, y_prob, pos_label = "Yes")

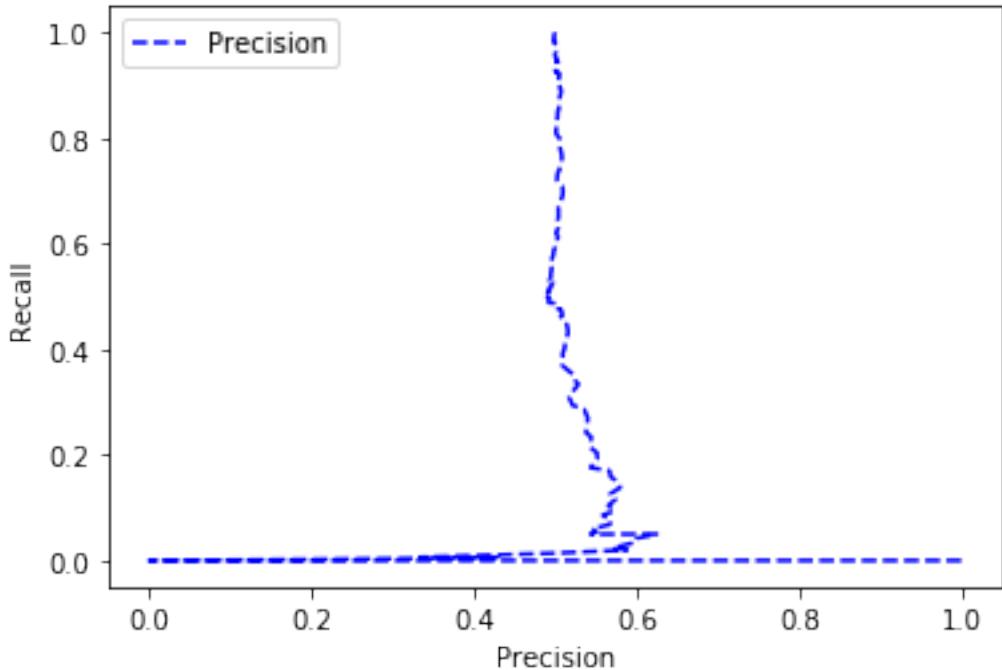
def plot_precision_recall_vs_threshold(precisions, recalls, thresholds):
    plt.plot(thresholds, precisions[:-1], "b--", label="Precision")
    plt.plot(thresholds, recalls[:-1], "g-", label="Recall")
    plt.xlabel("Threshold")
    plt.legend(loc="upper left")
    plt.ylim([0, 1])

plot_precision_recall_vs_threshold(precisions, recalls, thresholds)
plt.show()

plt.plot(precisions, recalls, "b--", label="Precision")
plt.xlabel("Precision")
plt.ylabel("Recall")
plt.legend(loc="upper left")
plt.show()

```

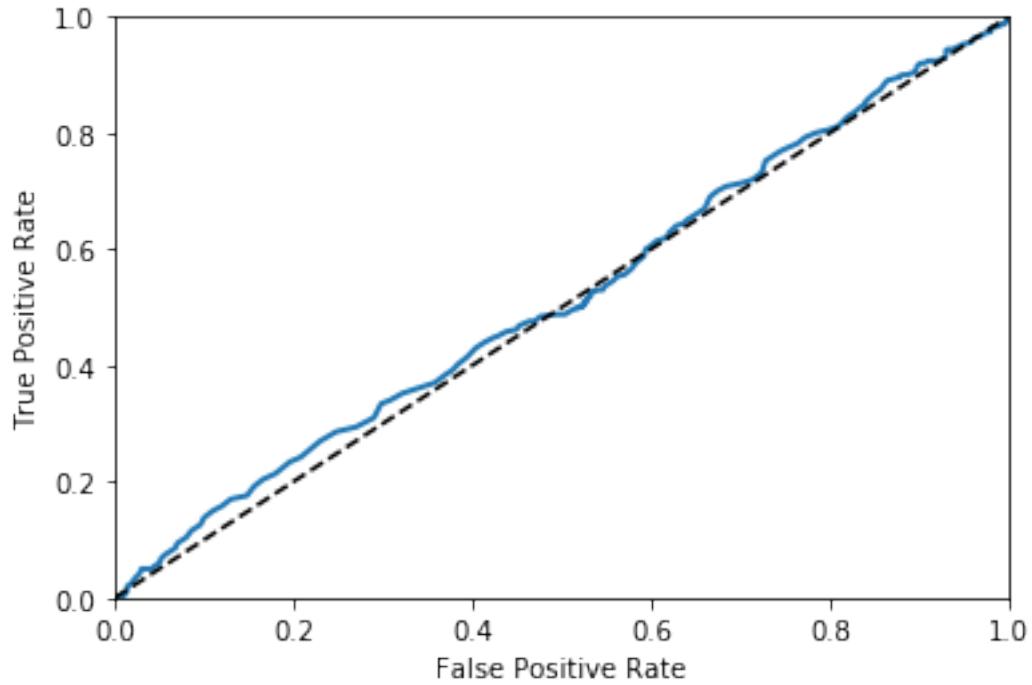




```
[17]: # ROC Curve
# The more the area under the curve the better our classifier
from sklearn.metrics import roc_curve
fpr, tpr, thresholds = roc_curve(y_test, y_prob, pos_label = "Yes")
def plot_roc_curve(fpr, tpr, label=None):
    plt.plot(fpr, tpr, linewidth=2, label=label)
    plt.plot([0, 1], [0, 1], 'k--')
    plt.axis([0, 1, 0, 1])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')

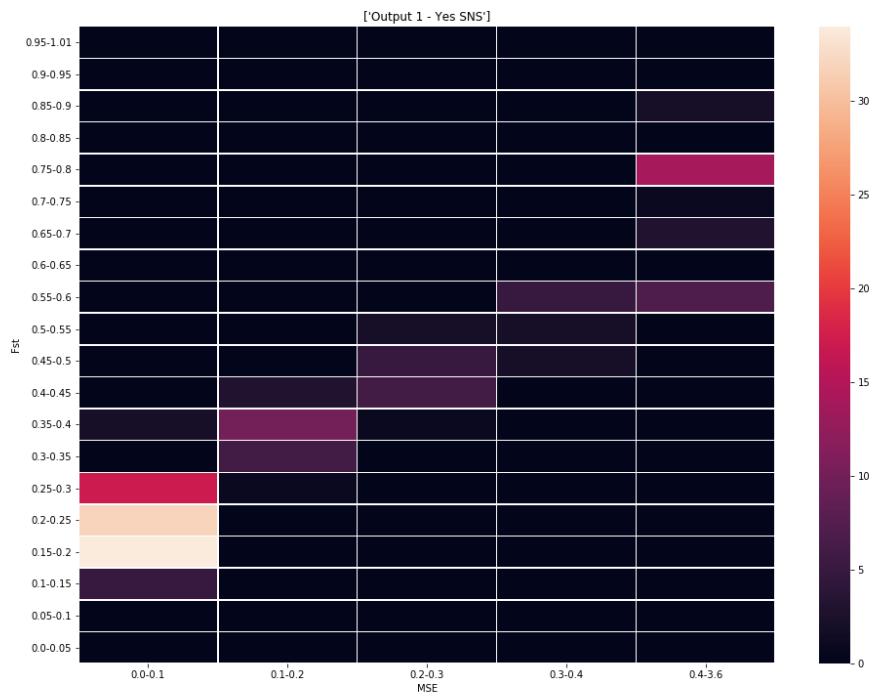
plot_roc_curve(fpr, tpr, "Random Forest")
plt.show()

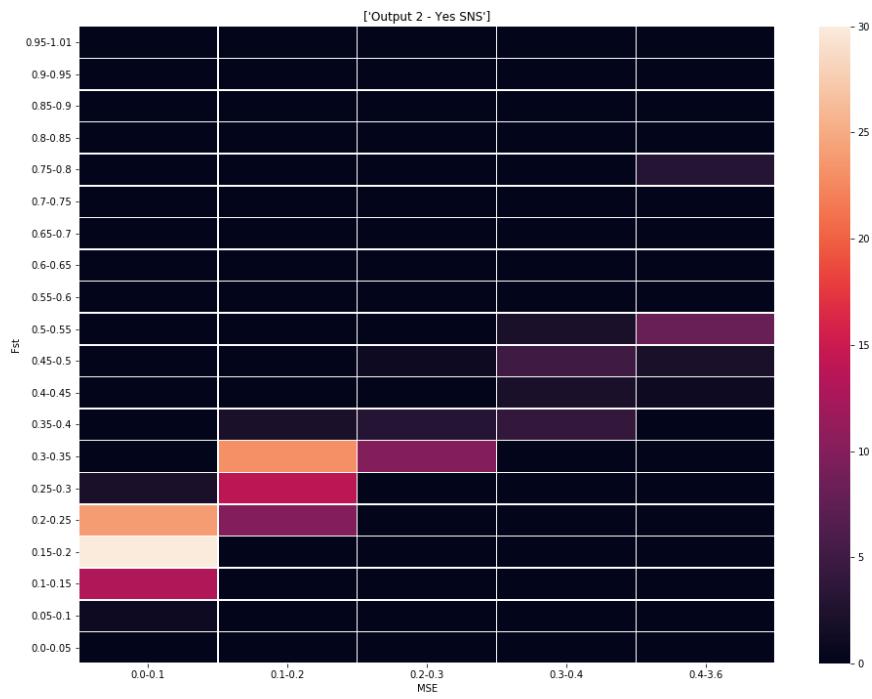
from sklearn.metrics import roc_auc_score
print("The AUC is ", roc_auc_score(y_test, y_prob))
```

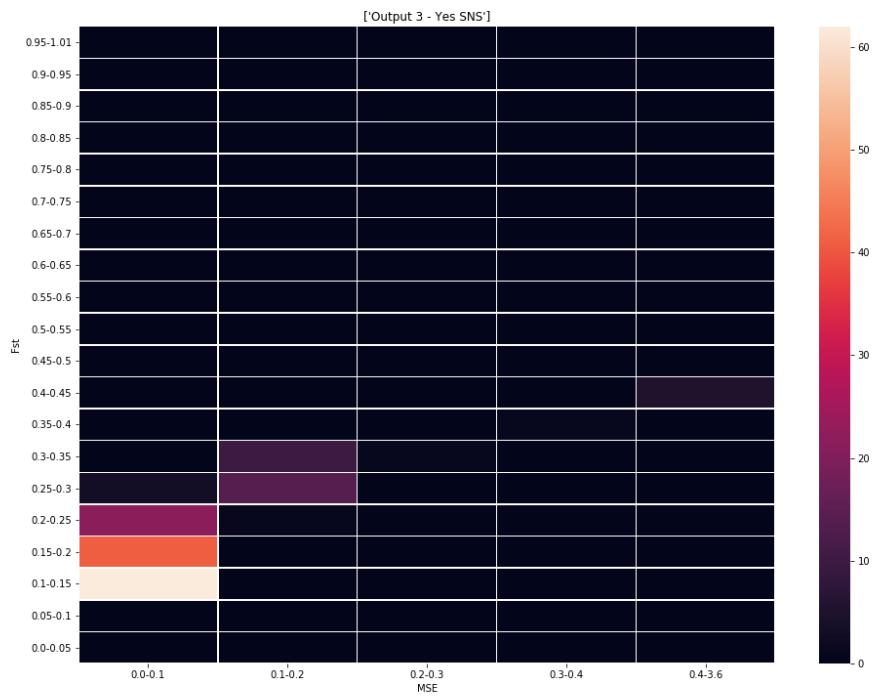


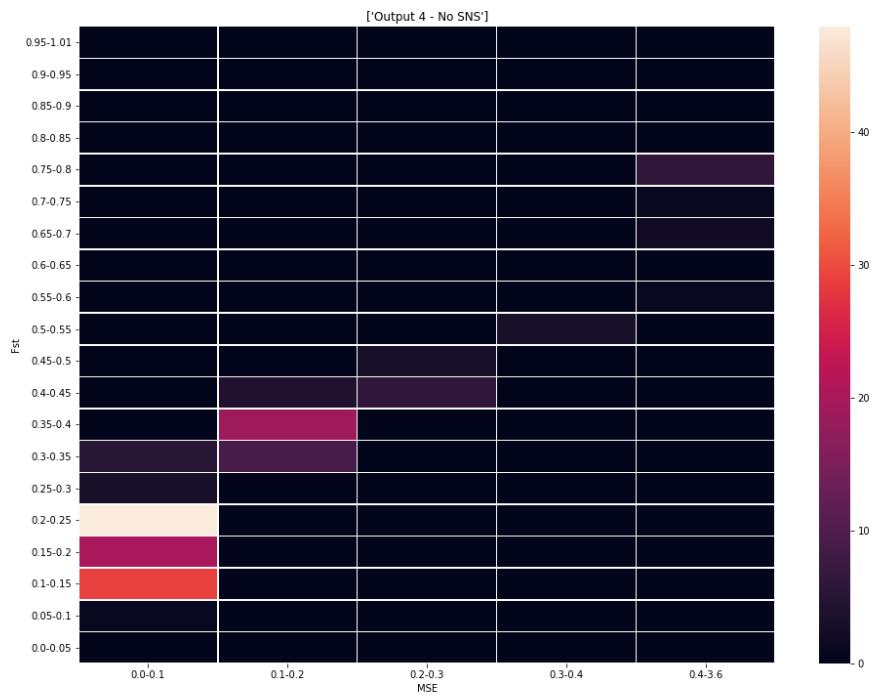
The AUC is 0.5153840000000001

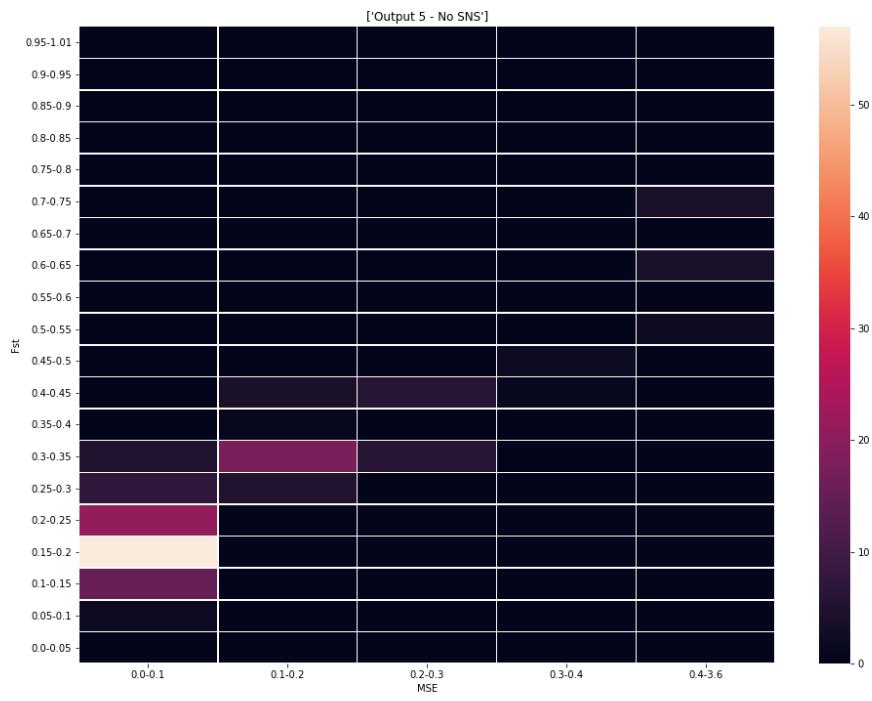
```
[18]: from IPython.display import Image
for i in range(1, 11):
    path = "./Images/output" + str(i) + ".png"
    display(Image(filename=path))
```

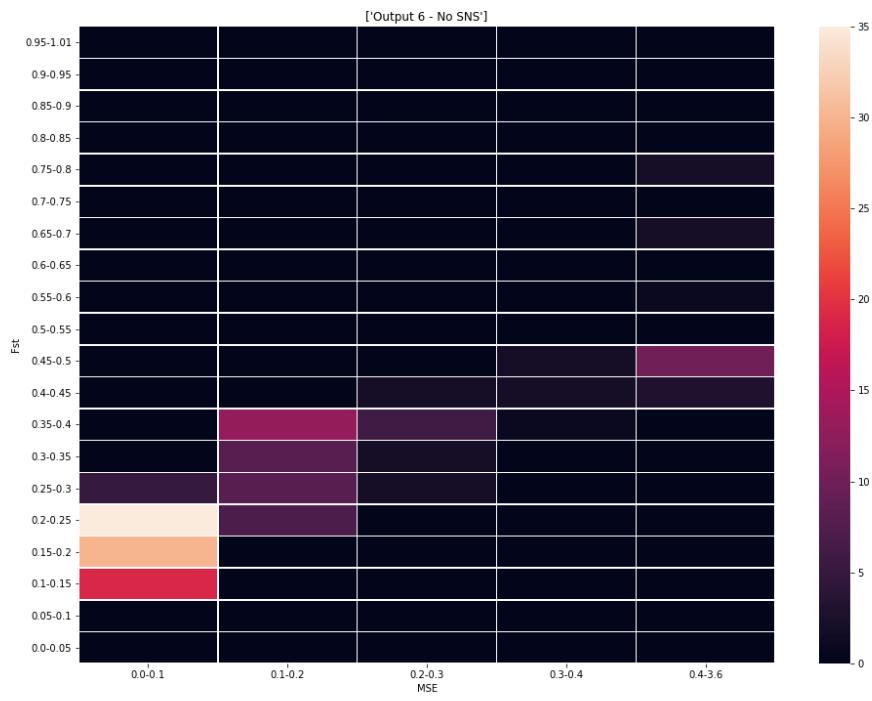


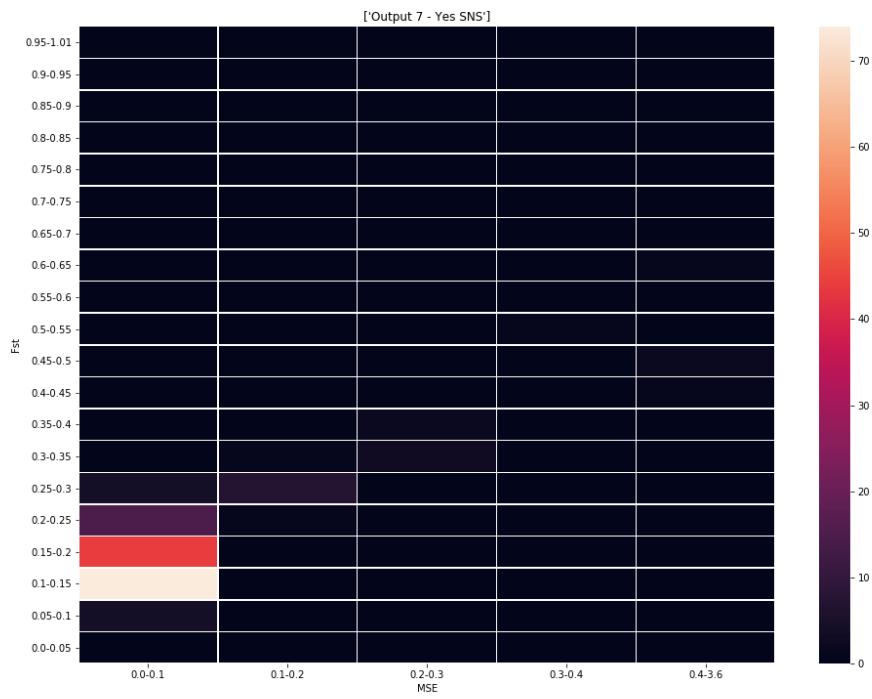


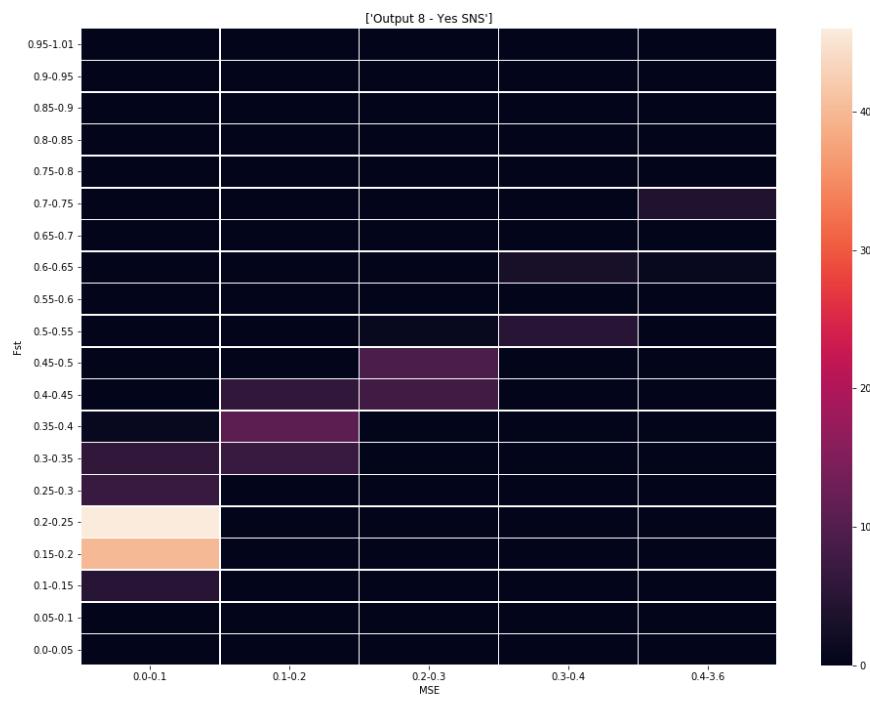


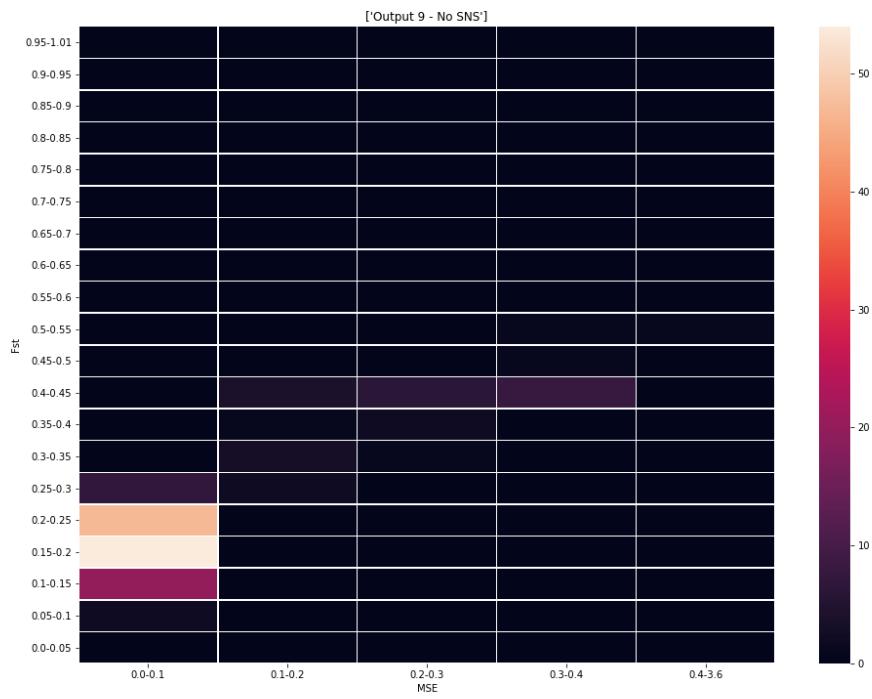


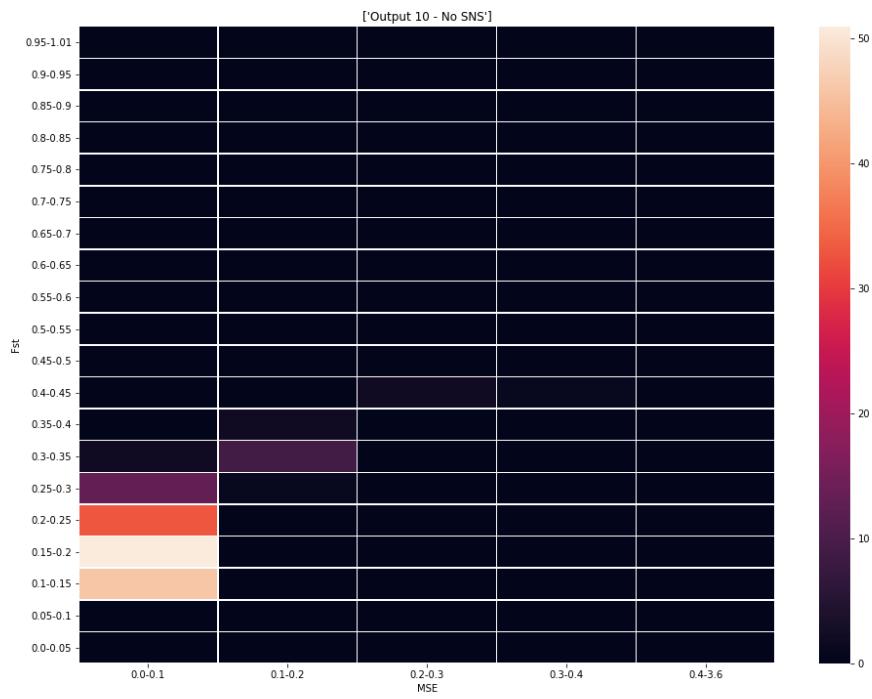












2 Issues

Can try a